

3-gVirtualXRay_vs_Gate-monochromatic-20keV

March 2, 2022

```
[1]: from IPython.display import display
from IPython.display import Image
from utils import * # Code shared across more than one notebook
```


Main contributors: F. P. Vidal and J. M. Létang

Purpose: In this notebook, we aim to demonstrate that gVirtualXRay is able to generate analytic simulations on GPU comparable to images generated with the state-of-the-art Monte Carlo simulation packages. We use here a monochromatic source of 20 keV.

Material and Methods: We simulate an image with gVirtualXRay and compare it with a ground truth image. For this purpose, we use [Gate](#), a wrapper for CERN's state-of-the-art Monte Carlo simulation tool: [Geant4](#). The number of tracked particles is 1e10.

In our simulation the source-to-object distance (SOD) is 1000mm, and the source-to-detector distance (SDD) is 1125mm. The beam spectrum is monochromatic, with an energy of 20 keV.


```
[2]: Image(filename="doc/setup.png")
```



The sample is made of a 70x70x15mm box of water, in which 5 columns of 3 spheres of different radii (2, 3.5, and 5mm) have been inserted. A given material is associated to the spheres of each column (bone (cortical), blood (whole), brain (grey/white matter), breast tissue, and adipose tissue). The columns are ordered in decreasing density. We use the definitions of tissue substitutes provided in the [ICRU Report 44](#) by the [International Commission on Radiation Units and Measurements](#). The material composition is available at <https://physics.nist.gov/PhysRefData/XrayMassCoef/tab2.html>.

```
[3]: Image(filename="doc/sample.png", width=400)
```

[3]:



Results: The calculations were performed on the following platform:

[4]: `printSystemInfo()`

OS:

Linux 5.3.18-150300.59.49-default
x86_64

CPU:

AMD Ryzen 7 3800XT 8-Core Processor

RAM:

63 GB

GPU:

Name: GeForce RTX 2080 Ti
Drivers: 455.45.01
Video memory: 11 GB

The Monte Carlo simulation needed 2.65e6 HS06 seconds to complete. It is equivalent to **1.15E+08**

ms (i.e. ~ 1.3 day) on the system used. Only 7 ± 2 ms was needed with the GPU used.

The **mean absolute percentage error (MAPE)**, also known as mean absolute percentage deviation (MAPD), between the two simulated images is **MAPE 0.69%**. The **zero-mean normalised cross-correlation** is **99.94%**. The **Structural Similarity Index (SSIM)** is **0.97**. As MAPE is low (close to 0), SSIM is high (close to 1), and ZNCC is high (close to 100%), we can conclude that this X-ray image simulated with gVirtualXRay on GPU in milliseconds is comparable to the same Monte Carlo simulation that ran for days.

1 Import packages

```
[5]: %matplotlib inline

import os # Locate files

import math
import numpy as np # Who does not use Numpy?
import pandas as pd # Load/Write CSV files

import matplotlib

from matplotlib.cm import get_cmap
import matplotlib.pyplot as plt # Plotting
from matplotlib.colors import LogNorm # Look up table
from matplotlib.colors import PowerNorm # Look up table
import matplotlib.colors as mcolors

font = {'family' : 'serif',
        #'weight' : 'bold',
        'size'   : 22
        }
matplotlib.rc('font', **font)
# matplotlib.rc('text', usetex=True)

from scipy.stats import pearsonr # Compute the correlatio coefficient

from skimage.util import compare_images # Checkboard comparison between two
↳ images
from skimage.metrics import structural_similarity as ssim
from sklearn.metrics import mean_absolute_percentage_error as mape
from skimage.metrics import structural_similarity as ssim

from tifffile import imread, imwrite # Load/Write TIFF files

import datetime # For the runtime
```

```
import viewscad # Use OpenSCAD to create STL files

import gvxrPython3 as gvxr # Simulate X-ray images

import json2gvxr # Set gVirtualXRay and the simulation up
```

SimpleGVXR 1.0.1 (2022-02-22T14:00:25) [Compiler: GNU g++] on Linux
gVirtualXRay core library (gvxr) 1.1.5 (2022-02-22T14:00:25) [Compiler: GNU g++]
on Linux

2 Reference image

We first load the reference image that has been simulated using [Gate](#) wrapper for CERN's [Geant4](#). Here we ignore scattering.

```
[6]: Image = imread("Gate_data/direct_20keV.tif") # Already corrected
Full_field = np.ones(Image.shape) # Perfect full field image
Dark_field = np.zeros(Full_field.shape) # Perfect dark field image
```

Projections are then corrected to account for variations in beam homogeneity and in the pixel-to-pixel sensitivity of the detector. This is the projection with flat-field correction (**Proj**):

$$\mathbf{Proj} = \frac{I - D}{F - D} \quad (1)$$

where F (full fields) and D (dark fields) are projection images without sample and acquired with and without the X-ray beam turned on respectively.

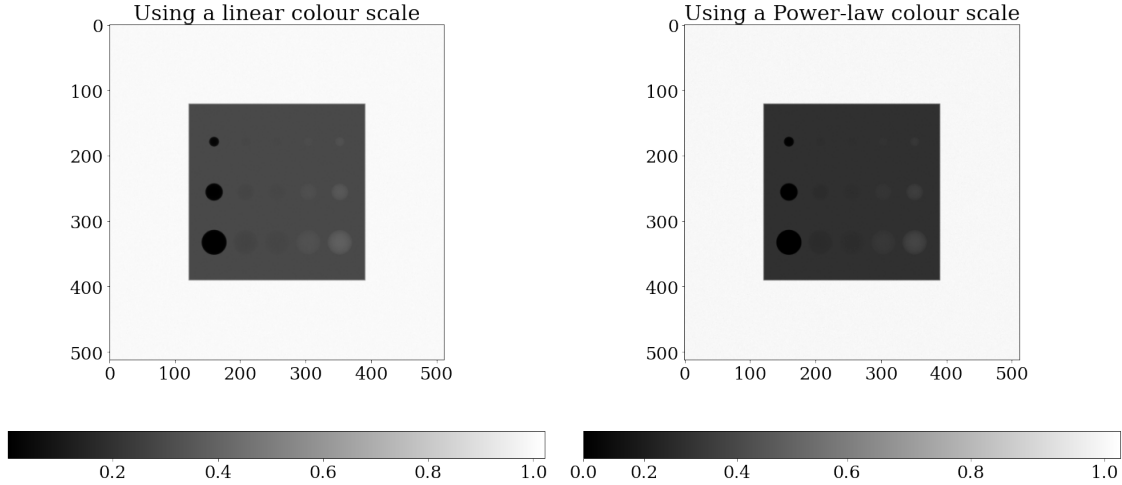
We now apply the flat-field correction to `Image`.

```
[7]: gate_image = (Image - Dark_field) / (Full_field - Dark_field)
# gate_image = Image / np.mean(Full_field)
```

We plot the image using a linear look-up table and a power-law normalisation.

```
[8]: displayLinearPowerScales(gate_image,
                               "Image simulated using Gate wrapper for CERN's Geant4",
                               ↪-- Mono-energy (20 keV)",
                               "plots/reference_from_Gate-monoE-20keV")
```

Image simulated using Gate wrapper for CERN's Geant4 -- Mono-energy (20 keV)



3 Setting up gVirtualXRay

Before simulating an X-ray image using gVirtualXRay, we must create an OpenGL context.

```
[9]: json2gvxr.initGVXR("notebook-3.json", "EGL")
```

Create an OpenGL context: 800x450

Wed Mar 2 12:33:45 2022 ---- Create window gvxrStatus: Create window

0

1.5

4.5.0 NVIDIA 455.45.01

Wed Mar 2 12:33:46 2022 ---- EGL version: Wed Mar 2 12:33:46 2022 ---- OpenGL

version supported by this platform OpenGL renderer: GeForce RTX 2080

Ti/PCIe/SSE2

OpenGL version: 4.5.0 NVIDIA 455.45.01

OpenGL vender: NVIDIA Corporation

Wed Mar 2 12:33:46 2022 ---- Use OpenGL 4.5.0 0 500 500

0 0 800 450

3.1 X-ray source

We create an X-ray source. It is a point source.

```
[10]: json2gvxr.initSourceGeometry()
```

Set up the beam

```
Source position: [0.0, 0.0, 1000.0, 'mm']  
Source shape: PointSource
```

3.2 Spectrum


The spectrum is monochromatic, with an energy of 20 keV.

```
[11]: spectrum, unit, k, f = json2gvxr.initSpectrum(verbose=1)  
      energy_set = sorted(spectrum.keys())  
  
      count_set = []  
  
      for energy in energy_set:  
          count_set.append(spectrum[energy])
```

1.0 photon of 20 keV

Plot the spectrum

```
[12]: plotSpectrum(k, f, 'plots/spectrum-monoE-20keV')
```



3.3 Detector

Create a digital detector

```
[13]: json2gvxr.initDetector()
```

Set up the detector

```
Detector position: [0.0, 0.0, -125.0, 'mm']
Detector up vector: [0, 1, 0]
Detector number of pixels: [512, 512]
Pixel spacing: [0.29296875, 0.29296875, 'mm']
```

3.4 Create and load the test object

We now create CAD models using [OpenSCAD](#) and extract the corresponding STL files.

```
[14]: openscad_make_spheres_str = """

module make_column_of(sphere_radius, height, count)
{
    step = height / (count - 1);
    for (a = [0 : count - 1]) {
        offset = -height / 2 + step * a ;
        translate([0, offset, 0])
            sphere(sphere_radius[a], $fn=25);
    }
}

module make_row_of(radius, count, id)
{
    step = radius / (count - 1);
    for (a = [0 : count - 1]) {
        if (id == -1 || id == a) {
            offset = -radius / 2 + step * a ;
            translate([offset, 0, 0])
                children();
        }
    }
}

module make_spheres(sphere_radius, ring_radius, ring_count, column_height,
    ↪column_count, id = -1)
{
    make_row_of(radius = ring_radius, count = ring_count, id = id)
        make_column_of(sphere_radius, height = column_height, count =
    ↪column_count);
}
"""
```

The matrix

```
[15]: openscad_matrix_str = """

color("red")
    difference() {
        scale([70, 70, 15])
            cube(1, center = true);
        make_spheres([2, 3.5, 5], 50, 5, 40, 3, -1);
    }

"""
```

```
[16]: fname = 'CAD_models/matrix.stl'
if not os.path.isfile(fname):

    r = viewscad.Renderer()
    r.render(openscad_matrix_str + openscad_make_spheres_str, outfile=fname)
```

```
[17]: openscad_cube_str = """

color("red")
    scale([70, 70, 15])
        cube(1, center = true);

"""
```

```
[18]: fname = 'CAD_models/cube.stl'
if not os.path.isfile(fname):

    r = viewscad.Renderer()
    r.render(openscad_cube_str, outfile='gvxr/input/cube.stl')
```

The spheres

```
[19]: openscad_col_str_set = []

for i in range(5):
    openscad_col_str_set.append("""
color("blue")
    make_spheres([2, 3.5, 5], 50, 5, 40, 3, "" + str(i) + ");")

    fname = 'CAD_models/col_' + str(i) + '.stl'
    if not os.path.isfile(fname):

        r = viewscad.Renderer()
        r.render(openscad_col_str_set[-1] + openscad_make_spheres_str,
        ↪outfile=fname)
```

Load the samples. `verbose=2` is used to print the material database for Gate. To disable it, use

verbose=0 or verbose=1.

```
[20]: json2gvxr.initSamples(verbose=1)
```

Load the 3D data

```
Load Bone_Cortical_ICRU_44 in CAD_models/col_0.stl using mm
Load Blood_Whole_ICRU_44 in CAD_models/col_1.stl using mm
Load Brain_Grey_White_Matter_ICRU_44 in CAD_models/col_2.stl using mm
Load Breast_Tissue_ICRU_44 in CAD_models/col_3.stl using mm
Load Adipose_Tissue_ICRU_44 in CAD_models/col_4.stl using mm
Load H2O in CAD_models/cube.stl using mm
```

```
CAD_models/col_0.stl    nb_faces:      1938    nb_vertices:    5814
bounding_box (in cm):  (-2.99606, -2.19961, -0.496354) (-2, 2.49901, 0.496354)
CAD_models/col_1.stl    nb_faces:      1938    nb_vertices:    5814
bounding_box (in cm):  (-1.74606, -2.19961, -0.496354) (-0.75, 2.49901,
0.496354)
CAD_models/col_2.stl    nb_faces:      1938    nb_vertices:    5814
bounding_box (in cm):  (-0.496057, -2.19961, -0.496354) (0.5, 2.49901,
0.496354)
CAD_models/col_3.stl    nb_faces:      1938    nb_vertices:    5814
bounding_box (in cm):  (0.753943, -2.19961, -0.496354) (1.75, 2.49901,
0.496354)
CAD_models/col_4.stl    nb_faces:      1938    nb_vertices:    5814
bounding_box (in cm):  (2.00394, -2.19961, -0.496354) (3, 2.49901, 0.496354)
CAD_models/cube.stl     nb_faces:       12     nb_vertices:     36
bounding_box (in cm):  (-3.5, -3.5, -0.75) (3.5, 3.5, 0.75)
```

4 Run the simulation

Update the 3D visualisation and take a screenshot

```
[21]: gvxr.displayScene()

gvxr.useLighting()
gvxr.useWireframe()
gvxr.setZoom(719.6787109375)
gvxr.setSceneRotationMatrix([0.7624880075454712, 0.09040657430887222, -0.
↪6406543850898743, 0.0,
                                0.05501500517129898, 0.9775413870811462, 0.
↪20342488586902618, 0.0,
                                0.6446591019630432, -0.190354123711586, 0.
↪7403913140296936, 0.0,
                                0.0, 0.0, 0.0, 1.0])

gvxr.displayScene()
```


```
[22]: screenshot = gvxr.takeScreenshot()
```

```
[23]: plt.figure(figsize= (10,10))
plt.title("Screenshot")
plt.imshow(screenshot)
plt.axis('off')

plt.tight_layout()

plt.savefig('plots/screenshot-beam-off-monoE-20keV.pdf')
plt.savefig('plots/screenshot-beam-off-monoE-20keV.png')
```

Screenshot



Compute an X-ray image 100 times (to gather performance statistics)

```
[24]: # gvxr.enableArtefactFilteringOnCPU()
gvxr.enableArtefactFilteringOnGPU()
# gvxr.disableArtefactFiltering() # Spere inserts are missing with GPU
↳ integration when a outer surface is used for the matrix

runtimes = []

for i in range(100):
    start_time = datetime.datetime.now()
    gvxr.computeXRayImage()
```

```

end_time = datetime.datetime.now()
delta_time = end_time - start_time
runtimes.append(delta_time.total_seconds() * 1000)

gvxr.displayScene()

```

```

[25]: screenshot = gvxr.takeScreenshot()


plt.figure(figsize= (10,10))
plt.title("Screenshot")
plt.imshow(screenshot)
plt.axis('off')

plt.tight_layout()

plt.savefig('plots/screenshot-beam-on-monoE-20keV.pdf')
plt.savefig('plots/screenshot-beam-on-monoE-20keV.png')

```

Screenshot



Save an X-ray image

```

[26]: # Compute the L-buffers on the GPU and integrate on the GPU
x_ray_image_integration_GPU = np.array(gvxr.computeXRayImage())
imwrite('gVirtualXRay_output_data/projection_raw_integration_GPU_monoE-20keV.
↳tif', x_ray_image_integration_GPU.astype(np.single))

```

Flat-field correction

```
[27]: total_energy_in_keV = 0.0
      for energy, count in zip(energy_set, count_set):
          total_energy_in_keV += energy * count

      total_energy_in_MeV = gvxr.getTotalEnergyWithDetectorResponse()
```

```
[28]: white = np.ones(x_ray_image_integration_GPU.shape) * total_energy_in_MeV
      dark = np.zeros(x_ray_image_integration_GPU.shape)

      x_ray_image_integration_GPU = (x_ray_image_integration_GPU - dark) / (white -
      ↪dark)
```

Save the corresponding image

```
[29]: imwrite('gVirtualXRay_output_data/
      ↪projection_corrected_integration_GPU_monoE-20keV.tif',
      ↪x_ray_image_integration_GPU.astype(np.single))
```

```
[30]: plt.figure(figsize= (20,10))

      plt.suptitle("Image simulated using gVirtualXRay,\nintegration on GPU", y=1.02)


      plt.subplot(121)
      plt.imshow(x_ray_image_integration_GPU, cmap="gray")
      plt.colorbar(orientation='horizontal')
      plt.title("Using a linear colour scale")

      plt.subplot(122)
      plt.imshow(x_ray_image_integration_GPU, norm=PowerNorm(gamma=1./0.75),
      ↪cmap="gray")
      plt.colorbar(orientation='horizontal')
      plt.title("Using a Power-law colour scale")

      plt.tight_layout()

      plt.savefig('plots/x_ray_image_integration_GPU-monoE-20keV.pdf')
      plt.savefig('plots/x_ray_image_integration_GPU-monoE-20keV.png')
```

Image simulated using gVirtualXRay,
integration on GPU



5 Comparison the analytic simulation with the Monte Carlo simulation

5.1 Quantitative validation

Compute image metrics between the two simulated images:

1. [mean absolute percentage error \(MAPE\)](#), also known as mean absolute percentage deviation (MAPD),
2. [zero-mean normalised cross-correlation \(ZNCC\)](#), and
3. [Structural Similarity Index \(SSIM\)](#).

We use these three metrics as one is a dissimilarity measurement (MAPE), two are similarity measurement (ZNCC & SSIM). MAPE and ZNCC can be expressed as a percentage, which eases the interpretation of the numerical values. SSIM is a number between 0 and 1. A good value of MAPE is 0%; of ZNCC 100%, and SSIM 1.

```
[31]: # Avoid div by 0
offset1 = min(gate_image.min(), x_ray_image_integration_GPU.min())
offset2 = 0.01 * (gate_image.max() - gate_image.min())
offset = offset2 - offset1

MAPE_integration_GPU = mape(gate_image + offset, x_ray_image_integration_GPU +
    ↪offset)
ZNCC_integration_GPU = np.mean((gate_image - gate_image.mean()) / gate_image.
    ↪std() * (x_ray_image_integration_GPU - x_ray_image_integration_GPU.mean()) /
    ↪x_ray_image_integration_GPU.std())
```

```
SSIM_integration_GPU = ssim(gate_image, x_ray_image_integration_GPU,
    ↪data_range=gate_image.max() - gate_image.min())

print("MAPE_integration_GPU:", "{0:0.2f}".format(100 * MAPE_integration_GPU) +
    ↪"%")
print("ZNCC_integration_GPU:", "{0:0.2f}".format(100 * ZNCC_integration_GPU) +
    ↪"%")
print("SSIM_integration_GPU:", "{0:0.2f}".format(SSIM_integration_GPU))
```

MAPE_integration_GPU: 0.69%
 ZNCC_integration_GPU: 99.94%
 SSIM_integration_GPU: 0.97

Get the total number of triangles

```
[32]: number_of_triangles = 0

for mesh in json2gvxr.params["Samples"]:
    label = mesh["Label"]
    number_of_triangles += gvxr.getNumberOfPrimitives(label)
```

```
[33]: runtime_avg = round(np.mean(runtimes))
runtime_std = round(np.std(runtimes))
```

Print a row of the table for the paper

```
[34]: print("Sphere inserts -- mono energy (20 keV) & Gate & " +
    "{0:0.2f}".format(100 * MAPE_integration_GPU) + "\\%    &    " +
    "{0:0.2f}".format(100 * ZNCC_integration_GPU) + "\\%    &    " +
    "{0:0.2f}".format(SSIM_integration_GPU) + "    &    $" +
    str(json2gvxr.params["Detector"]["NumberOfPixels"][0]) + " \\pm " +
    ↪str(json2gvxr.params["Detector"]["NumberOfPixels"][1]) + "$    &    " +
    str(number_of_triangles) + "    &    " +
    "1.15E+08    &    " +
    "$" + str(runtime_avg) + " \\pm " + str(runtime_std) + "$ \\\\" )
```

Sphere inserts -- mono energy (20 keV) & Gate & 0.69% & 99.94% &
 0.97 & \$512 \pm 512\$ & 9702 & 1.15E+08 & \$8 \pm 2\$ \

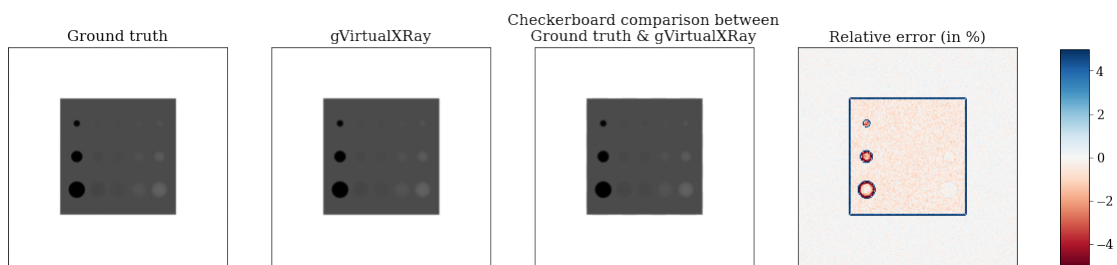
In both cases, MAPE is very small (less than 1%), ZNCC is very high (more than 99%), and SSIM is very high (almost 1). We can conclude that the two images are similar. The main difference lie in the Poisson noise affecting the Monte Carlo simulation.

5.2 Qualitative validation

Checkboard comparison

```
[35]: font = {'size' : 12.5
           }
matplotlib.rc('font', **font)
```


```
[36]: fullCompareImages(gate_image=gate_image,
                        gvxr_image=x_ray_image_integration_GPU,
                        title="gVirtualXRay",
                        fname="plots/full_comparison_integration_GPU-monoE-20keV",
                        log=False,
                        vmin=0.0,
                        vmax=1,
                        avoid_div_0=True)
```




Plot the profiles

```
[37]: font = {'size' : 22
           }
matplotlib.rc('font', **font)
```

```
[38]: plotTwoProfiles(json2gvxr, gate_image, x_ray_image_integration_GPU, "plots/
    ↪profiles-checkerboard-monoE-20keV")
```



```
[39]: spacing = json2gvxr.params["Detector"]["Size"][0] / json2gvxr.
      ↪params["Detector"]["NumberOfPixels"][0]
min_limit = round(40)
max_limit = round(512 * spacing - 40)
plotTwoProfiles(json2gvxr, gate_image, x_ray_image_integration_GPU, "plots/
      ↪profiles-zoom-checkerboard-monoE-20keV", [min_limit, max_limit])
```

6 All done

Destroy the window

```
[40]: gvxr.destroyAllWindows()
```

0(0x563752cfddc0)